

**AutoIdeaFlow: from Idea Generation to
Paper Writeup and Review**

A Minor Project Report

Submitted To



Chhattisgarh Swami Vivekanand Technical University

Bhilai

For

Minor Project

of

Bachelor of Technology (Hons.)

in

Computer Science & Engineering

By

Jayant Patel

300012721061

CB4646

7th Sem

Artificial Intelligence

Under the Guidance of

Dr. Nachiket Tapas

Assistant Professor

Department of Computer Science & Engineering

UTD, CSVTU, Bhilai (C.G.)



**Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107**

Session: 2024 – 25



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

DECLARATION BY THE CANDIDATE

We the undersigned solemnly declare that the Minor project report entitled “***AUTOIDEAFLOW FROM IDEA GENERATION TO PAPER WRITEUP AND REVIEW***” is based our own work carried out during the course of our study under the supervision of ***Dr. Nachiket Tapas*** and ***Mr. Abhinaw Jagtap***.

We assert that the statements made and conclusions drawn are an outcome of the project work. We further declare that to the best of our knowledge and belief that the report does not contain any part of any work which has been submitted for the award of any other degree/diploma/certificate in this University/Deemed university of India or any other country.

Jayant Patel

Roll No: 300012721061

Enroll No: CB4646

Semester: 7th (CSE)



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

CERTIFICATE BY THE SUPERVISOR

This is to certify that the Minor project report entitled “*AUTOIDEAFLOW FROM IDEA GENERATION TO PAPER WRITEUP AND REVIEW*” is a record of project work carried out under my guidance and supervision for the fulfillment of the award of degree of Bachelor of Technology (Hons.) in the faculty of Computer Science & Engineering of Chhattisgarh Swami Vivekananda Technical University, Bhilai (C.G.) India.

To the best of my knowledge and belief the report

- i. Embodies the work of the candidate himself.
- ii. Has duly been completed.
- iii. Fulfills the partial requirement of the ordinance relating to the B.Tech. (Hons) degree of the University.
- iv. Is up to the desired standard both in respect of contents and language for being referred to the examiners.

Dr. Nachiket Tapas

Assistant Professor

Depr. of Computer Science &
Engineering, UTD, CSVTU, Bhilai
(C.G.)

Forwarded to
Chhattisgarh Swami Vivekanand Technical University, Bhilai (C.G.)

Dr. J P Patra

HOD

Dept. of Computer Science &
Engineering, UTD, CSVTU, Bhilai
(C.G.)

Dr. P K Ghosh

Director

UTD, CSVTU, Bhilai (C.G.)



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

CERTIFICATE BY THE EXAMINERS

The project report entitled “*AUTOIDEAFLOW FROM IDEA GENERATION TO PAPER WRITEUP AND REVIEW*” has been examined by the undersigned as a part of the examination of Bachelor of Technology (Hons.) in the faculty of Computer Science & Engineering of Chhattisgarh Swami Vivekanand Technical University, Bhilai.

Internal Examiner
Date:

External Examiner
Date:



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

ACKNOWLEDGEMENT

Working for this project has been a great experience for us. There were moments of anxiety, when we could not solve a problem for the several days. But we have enjoyed every bit of process and are thankful to all people associated with us during this period we convey our sincere thanks to our project guide **Dr. Nachiket Tapas** and co-guide **Mr. Abhinaw Jagtap** for providing me all sorts of facilities. His support and guidance helped us to carry out the project. We owe a great dept. of his gratitude for his constant advice, support, cooperation & encouragement throughout the project we would also like to express our deep gratitude to respected **Dr. J P Patra** (Head of Department) for his ever helping and support. We also pay special thanks for his helpful solution and comments enriched by his experience, which improved our ideas for betterment of the project. We would also like to express our deep gratitude to respected **Dr. P K Ghosh** (Director) and college management for providing an educational ambience. It will be our pleasure to acknowledge, utmost cooperation and valuable suggestions from time to time given by our staff members of our department, to whom we owe our entire computer knowledge and also we would like to thank all those persons who have directly or indirectly helped us by providing books and computer peripherals and other necessary amenities which helped us in the development of this project which would otherwise have not been possible.

Jayant Patel

Roll No: 300012721061

Enroll No: CB4646

Semester: 7th (CSE (AI))



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

TABLE OF CONTENTS

Declaration by the Candidate	I
Certificate by the Supervisor	II
Certificate by the Examiners	III
Acknowledgement	IV
Table of Contents	V
List of Abbreviations	VII
List of Figures	VIII
Abstract	IX
1 Introduction	1
1.1 Background Information	1
1.1.1 Large Language Models	2
1.1.2 LLM Agent Framework	3
1.1.3 Aider: An LLM-Bases Coding Assistant	3
1.2 Aims and Objectives	4
1.3 Significance of the Project	5
1.4 Scope and Limitations	7
1.4.1 Scope	7
1.4.2 Limitations	7
1.5 Software Engineering Paradigms	9
1.5.1 Procedural Programming Paradigm	9
1.5.2 Object-Oriented Programming (OOP) Paradigm	10
1.5.3 Functional Programming Paradigm	10
1.5.4 Agile Software Development Paradigm	10
1.5.5 DevOps Paradigm	11
1.5.6 Software as a Service (SaaS) Paradigm	11
1.6 Overview of the Structure	11
2 Methodology	14
2.1 Project Overview	14
2.2 Research Design and Approach	15
2.3 Data Collection Methods	17
2.4 Project Directory	18
2.4.1 Structure	18
2.4.2 Explanation of Structure	18
2.5 Data Analysis Techniques	19
2.6 Ethical Considerations	20
2.7 Limitations	20
3 Implementation	22
3.1 Development Environment	22

V

3.1.1	Software Tools and Frameworks	22
3.1.2	Hardware Resources	23
3.1.3	Collaboration Tools	23
3.2	Project Implementation	24
3.2.1	Execution Stages	24
3.3	Project Timeline	24
3.4	Resource Management	25
3.4.1	Human Resources	25
3.4.2	Computational Resources	26
3.4.3	Financial Resources	26
3.5	Challenges Faced	27
3.6	Lessons Learned	28
4	Results and Discussion	29
4.1	Presentation of Results	29
4.1.1	Visual Analysis	29
4.1.2	Statistical Analysis	30
4.1.3	Comparative Analysis	31
4.2	Interpretation of Results	32
4.2.1	General Observations	32
4.2.2	Comparisons with Established Models	33
4.2.3	Addressing Research Challenges	34
4.2.4	Review Generated by our Framework	34
4.3	Discussion	37
5	Conclusion	39
5.1	Achivement of the Objectives	39
5.2	Implications and Recommendations	39
5.2.1	General Implications	39
5.2.2	Key Recommendations	40
5.3	Future Scope	40
5.3.1	Development Paths	41
5.3.2	Additional Future Directions	42
5.4	Personal Reflection	43
5.5	Summary of the key findings	43
5.5.1	Key Points	43
5.5.2	Final Remarks	44
	References	45



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

LIST OF ABBREVIATIONS

Abbreviation	Description
AI	Artificial Intelligence
ML	Machine Learning
LLM	Large Language Model
DDPM	Denoising Diffusion Probabilistic Model
VAE	Variational Autoencoder
GAN	Generative Adversarial Network
MLP	Multi-Layer Perceptron
RMSE	Root Mean Squared Error
MSE	Mean Squared Error
KPI	Key Performance Indicator
GPU	Graphics Processing Unit
CPU	Central Processing Unit
API	Application Programming Interface



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

LIST OF FIGURES

1.1	Process Flow Diagram Highlighting Bottlenecks in Traditional Research Proposal Methods	6
2.1	High-level overview of the research generation workflow.	14
2.2	Main modules of the automated research generation system.	15
2.3	Detailed system processes for automated research generation.	16
4.1	Plots generated by the framework showing training loss over epochs for the generated paper's model.	29
4.2	Visualization of generated output images from the framework	30
4.3	Title page of the research paper generated by our framework.	31
4.4	Result page of the research paper generated by our framework.	32
4.5	References generated by our framework using Semantic Scholar API.	33



Department of Computer Science & Engineering
University Teaching Department
Chhattisgarh Swami Vivekanand Technical University
Bhilai (C.G.) 491107

ABSTRACT

Developing agents capable of autonomously conducting scientific research and discovering new knowledge is a major challenge in artificial general intelligence. While large language models (LLMs) have assisted human scientists in tasks like brainstorming and coding, they handle only a fraction of the scientific process. This paper introduces The AutoIdeaFlow, the first comprehensive framework enabling LLMs to perform end-to-end scientific discovery, including idea generation, coding, experimentation, result visualization, paper writing, and simulated review. Applied to three machine learning subfields—diffusion modeling, transformer-based language modeling, and learning dynamics—The AutoIdeaFlow can develop full papers at under \$15 each, democratizing research and accelerating scientific progress. An automated reviewer evaluates these papers with near-human accuracy, demonstrating that The AutoIdeaFlow can produce work meeting top conference standards. This framework heralds a new era of AI-driven scientific discovery, fostering continuous innovation on complex global challenges.

Keywords: Automated Research, Scientific Discovery, Large Language Models, Experiment Automation, Novelty Assessment, Paper Generation, Semantic Scholar API, Automated Peer Review.

Chapter 1

Introduction

1.1 Background Information

Traditionally, scientific discovery has always been a time-consuming and trial-by-error process that is guided by the systematic steps, which a researcher follows in order to improve knowledge. This pattern usually involves mapping of the unknown, generating hypotheses, conducting experiments, analyzing results, and communication of the findings. As much as this comprehensive process has produced significant advancement across many disciplines, it inherently suffers from human limitations regarding time, creativity, and availability of resources. With the rising need for more effective research methodologies, there is a growing interest in utilizing automation and computational techniques to improve the research process.

The latest developments in computational techniques and machine learning have made it possible to automate the most diverse components of scientific inquiry. The modern automatic systems can assist researchers with literature reviews, data analysis, and designing experiments. Such systems use advanced computational models that can understand and generate natural language, coding, and compiling reports. Such developments have the potential to greatly accelerate the pace of research and make scientific information more available by reducing the costs and efforts required to produce high-quality output. Despite these advances, the full automation of scientific inquiry across its entire life cycle is still a far-from-reachable goal. Current systems are often confined to a specific domain or task and typically require substantial human intervention. For instance, while the automated equipment has the capacity to conduct experiments on its own, it is the human researchers who still control which experiments should be performed.

To address these challenges, there is a critical need for comprehensive frameworks [4] that can manage the entire research process—including idea generation, experimental execution, and documentation—completely autonomously. In this project, we explore and implement such a framework, testing its effectiveness in multiple research domains and exploring possibilities for improvement. By including various automation technologies in scientific workflows, we will look to optimize processes, increase reproducibility, and encourage collaborative research efforts.

This has significant consequences. As automation systems become more complex, they may enable researchers to focus on higher-level conceptual tasks while delegating routine analyses and experi-

mental procedures to machines. This shift is potentially capable of leading to discoveries at a faster pace and a more inclusive scientific community in which access to high-end research instruments is equitably made.

Advancements in machine learning [2], particularly in the development of large language models (LLMs) and other AI technologies, have opened new avenues for automating various aspects of scientific research. These advancements enable the creation of intelligent systems that can perform complex tasks such as data analysis, hypothesis generation, and even experimental design with minimal human intervention. AI-driven automation can significantly enhance the efficiency and accuracy of research processes by leveraging vast amounts of data and sophisticated algorithms to uncover patterns and insights that might be missed by human researchers.

In the context of this project, AI can be utilized to automate several key components. For instance, machine learning algorithms can be employed to analyze large datasets, identify trends, and generate hypotheses based on the observed data. Natural language processing (NLP) techniques can assist in literature reviews by automatically summarizing relevant research papers and extracting key information. Additionally, AI can be used to design and optimize experiments, ensuring that they are conducted in the most efficient and effective manner possible. By integrating these AI capabilities into the research framework, we aim to create a system that not only accelerates the research process but also improves the quality and reproducibility of scientific findings.

1.1.1 Large Language Models

Large Language Models (LLMs) are advanced machine learning systems trained to process and generate human-like text based on a given prompt [1] & [3]. They are typically built using transformer architectures, which excel at capturing contextual relationships in sequential data. LLMs are trained on vast amounts of text data to learn statistical patterns, enabling them to perform various tasks such as text generation, translation, summarization, and question-answering.

The core functionality of an LLM revolves around predicting the probability of the next word or token in a sequence, conditioned on the preceding context. This allows LLMs to generate coherent and contextually relevant outputs. Over time, LLMs like GPT-3, GPT-4, [5] and others have demonstrated impressive capabilities, including reasoning, coding, and creating content that appears human-authored.

LLMs have been successfully applied in diverse domains, including but not limited to:

- Natural Language Processing (NLP): Text completion, summarization, and sentiment analysis.
- Scientific Research: Generating hypotheses, writing papers, and assisting in literature reviews.

- **Code Generation:** Helping developers write, debug, and optimize code.
- **Content Creation:** Crafting articles, reports, and other creative works.

1.1.2 LLM Agent Framework

Large Language Models (LLMs) are advanced machine learning systems trained to process and generate human-like text based on a given prompt. They are typically built using transformer architectures, which excel at capturing contextual relationships in sequential data. LLMs are trained on vast amounts of text data to learn statistical patterns, enabling them to perform various tasks such as text generation, translation, summarization, and question-answering.

The core functionality of an LLM revolves around predicting the probability of the next word or token in a sequence, conditioned on the preceding context. This allows LLMs to generate coherent and contextually relevant outputs. Over time, LLMs like GPT-3, GPT-4, and others have demonstrated impressive capabilities, including reasoning, coding, and creating content that appears human-authored.

LLMs have been successfully applied in diverse domains, including but not limited to:

1. **Natural Language Processing (NLP):** Text completion, summarization, and sentiment analysis.
2. **Scientific Research:** Generating hypotheses, writing papers, and assisting in literature reviews.
3. **Code Generation:** Helping developers write, debug, and optimize code.
4. **Content Creation:** Crafting articles, reports, and other creative works.

1.1.3 Aider: An LLM-Bases Coding Assistant

Aider is an open-source coding assistant designed to automate and streamline the software development process. It uses the capabilities of LLMs to understand natural language instructions, perform code generation, fix bugs, refactor existing codebases, and even implement new features based on developer input.

Key Features of Aider:

- **Code Implementation:** Aider can understand the context of existing codebases and add new functionalities based on user prompts.
- **Error Handling:** It identifies bugs and suggests fixes, enabling developers to debug their code more efficiently.

- **Refactoring:** Aider can improve code readability, structure, and maintainability through automatic refactoring.
- **Advanced Integration:** It can seamlessly integrate with various software libraries and tools, making it suitable for complex coding tasks.

Aider leverages cutting-edge LLM capabilities to achieve high success rates in implementing requested changes. For instance, its reliability has been benchmarked at approximately 18.9% success on the SWE Bench, a collection of real-world GitHub issues.

1.2 Aims and Objectives

The main objective of this system is to create a completely autonomous form that can transform the existing research process by automating the critical components of scientific research. The system, with an integration of advanced computational techniques and machine learning, along with automated tools, aims to enhance the efficiency, reproducibility, and accessibility of scientific research. The overall objective remains to streamline the research process while allowing for faster discoveries and fostering collaboration, thereby democratizing access to high-quality research tools and methodologies.

To achieve this aim, the project is guided by the following objectives:

- **Idea Generation:** The system will leverage computational creativity and advanced algorithms to autonomously generate innovative research ideas. These ideas will be novel, feasible, and aligned with the specific needs and challenges of a given domain. The objective is to inspire groundbreaking research directions by harnessing the power of automation to explore a vast landscape of possibilities beyond human capabilities.
- **Experimental Design and Execution:** A sound framework will be built for designing and executing experiments or simulations without human intervention. This framework will include choosing the experimental parameters, setting up experimental setups, and result collection. It will optimize experimental workflows in order to minimize human intervention and increase efficiency, accuracy, and scalability of research.
- **Result Analysis and Logging:** The system will systematically analyze experimental outcomes based on advanced data analysis techniques. All findings will be logged, interpreted, and presented in a logical, structured, and academically compliant format so that results are not only

accurate and meaningful but also effectively communicated in a way that reflects the highest scientific standards.

- **Peer Review Simulation:** To ensure the quality and credibility of the research produced, the system will have a simulation of the peer-review process. This feature will check the relevance, originality, and rigour of the research outcomes. It will provide constructive feedback for maintaining high standards and helping in iterative improvement.
- **Cost-Effectiveness:** The project recognizes the importance of accessibility and emphasizes that the system should operate within reasonable computational and financial parameters. This objective ensures that the system remains accessible to researchers and organizations with limited resources, thereby allowing broader participation and fostering inclusivity in scientific discovery.

1.3 Significance of the Project

The significance of this project lies in its potential to transform how research is conducted and disseminated. Traditional research workflows often require substantial time, expertise, and resources, which can limit participation to well-funded institutions or individuals with specialized skills. This project addresses these challenges by focusing on several key areas: The significance of this project lies in its potential to transform how research is conducted and disseminated. Traditional research workflows often require substantial time, expertise, and resources, which can limit participation to well-funded institutions or individuals with specialized skills. This project addresses these challenges by focusing on several key areas.

Firstly, by automating the research process, the project aims to reduce barriers for individuals or organizations with limited resources. This enhancement of accessibility will enable broader participation in scientific inquiry, allowing more diverse voices and perspectives to contribute to research efforts. Secondly, the implementation of streamlined workflows and the elimination of bottlenecks are central to this project. By doing so, the system will facilitate faster hypothesis testing and knowledge generation, significantly increasing the pace at which new discoveries can be made. Thirdly, automated systems can standardize experimental procedures and documentation, thereby reducing human errors and improving the reproducibility of research findings. This improvement is crucial for maintaining the integrity of scientific research and ensuring that results can be reliably replicated by other researchers.

Additionally, by reducing the cost of conducting and publishing research, this project empowers

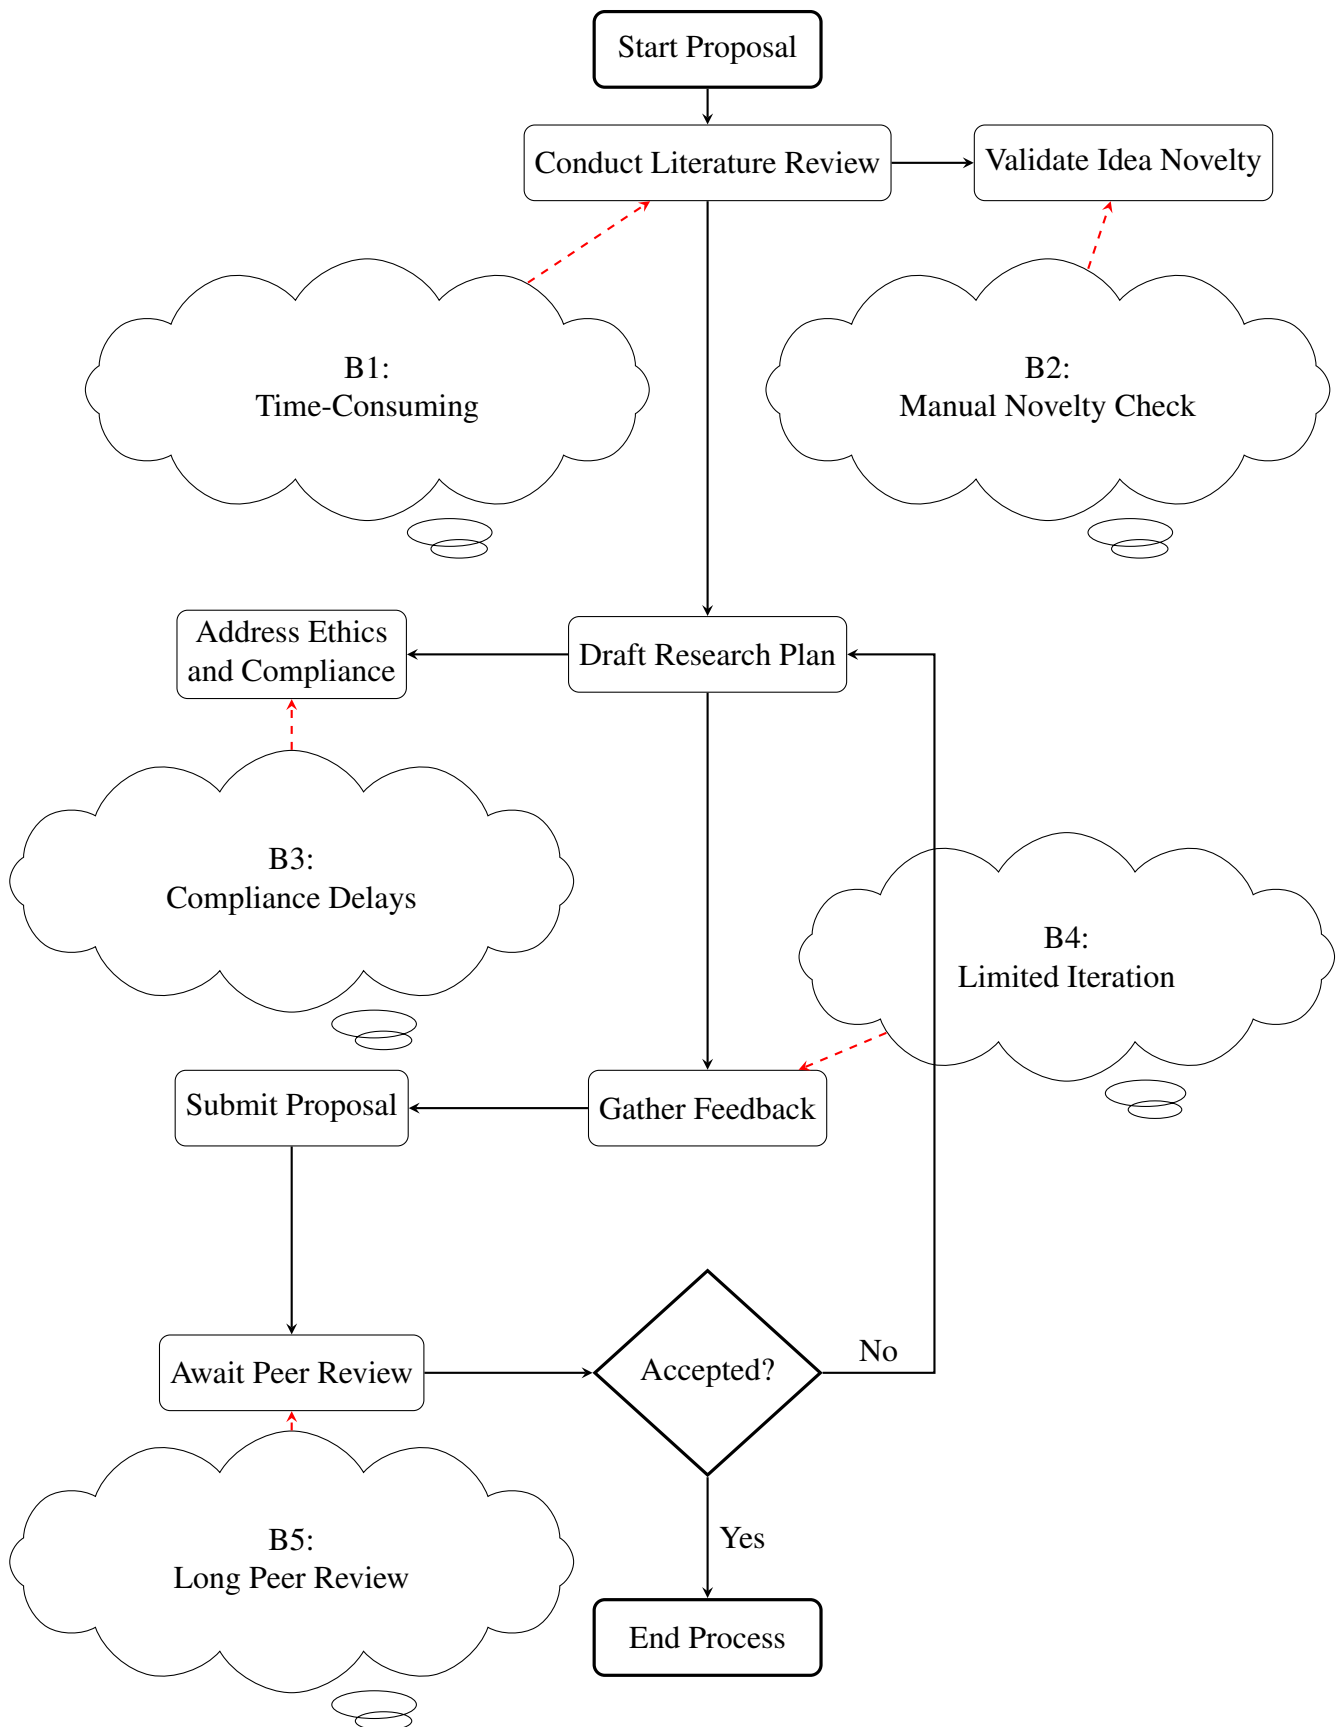


Figure 1.1: Process Flow Diagram Highlighting Bottlenecks in Traditional Research Proposal Methods

smaller teams and underrepresented regions to contribute to global scientific advancements. This democratization of innovation fosters a more inclusive scientific community where diverse ideas can flourish. Finally, automation has the potential to encourage cross-domain exploration by minimizing the need for domain-specific expertise during the initial stages of hypothesis generation and experimentation. This capability can lead to novel interdisciplinary collaborations that might not have been possible within traditional research frameworks.

1.4 Scope and Limitations

1.4.1 Scope

It is this initiative which aims to automate the basic parts of the research process especially with regard to idea generation experimentation, and result-documentation. While it's basically tested within a domain, namely computational science or machine learning, the base framework is quite easily modified for application to other disciplines with suitable adaptation. The initiative further incorporates several tools and methodologies dedicated to the assessment of novelty, ensuring that the ideas produced do not simply replicate existing works. Through the implementation of automated review processes, the initiative aspires to emulate peer-review criteria while offering a thorough evaluation of the quality of research. This comprehensive strategy intends to improve both the integrity and significance of the research outputs produced by the system.

1.4.2 Limitations

Although the scope of this project is vast, it has to be well noted that some intrinsic boundaries exist:

1. **Domain-specific limitations:** The system is likely to face challenges in specific domains that require exclusive information or proprietary datasets. It may not be so effective in certain research domains in which a subtle understanding is required.
2. **Dependency on Computation Resources:** Availability and cost for computation resources determine the running cost of the system directly. Change in resource availability will have impact on the use of functionality by different types of users or organizations.
3. **Implementation Challenges:** There is a possibility that bugs in the implementation have generated misleading results or part analyses due to the limitations. Accuracy and reliability are major challenges, and the algorithms at present need to be continuously refined.

4. **Ethical Concerns:** There may be an opportunity for its abuse, such as preparing pseudo-scientific reports or unethical research submissions. These issues highlight the necessity of constant supervision and regulation in minimizing the risks involved with the automated production of research.
5. **Human Oversight Required:** It does produce results and reports with results, but the wider implication often requires human insight, which limits the complete automation of the system and therefore underlines the importance of cooperation between automated tools and human researchers.
6. **Current Failure Modes** The framework, in its current form, has several shortcomings in addition to those already identified. These include, but are not limited to:
 - The idea generation process often results in very similar ideas across different runs and models. This issue may be addressed by allowing the system to follow up and delve deeper into its best ideas or by providing it with content from recently published projects as a source of novelty.
 - There is a failure to implement a significant fraction of the proposed ideas. Additionally, there are frequent issues with generating LaTeX that compiles correctly. While the system can produce creative and promising ideas, many are too challenging for it to implement effectively.
 - The framework may incorrectly implement an idea, which can be difficult to catch. An adversarial code-checking reviewer may partially address this issue; however, manual verification of implementations is essential before trusting reported results.
 - Due to the limited number of experiments conducted per idea, the results often do not meet the expected rigor and depth of a standard machine learning conference project. Moreover, the constraints on the number of experiments hinder fair comparisons that control for parameters, FLOPs, or runtime, leading to potentially deceptive or inaccurate conclusions. These issues are expected to improve as the costs of compute and foundation models decrease.
 - Currently, without utilizing vision capabilities, the system cannot correct visual issues in its outputs or interpret plots. For instance, generated plots may be unreadable, tables may exceed page width, and overall layout quality is often suboptimal. Future versions with integrated vision capabilities should address these concerns.

- When writing, the framework sometimes struggles to find and cite the most relevant projects. It also frequently fails to reference figures correctly in LaTeX and may hallucinate invalid file paths.
- Importantly, critical errors can occur when writing and evaluating results. For example, it struggles with comparing magnitudes of numbers—a known issue with LLMs. Additionally, when changing metrics (e.g., loss functions), it sometimes fails to consider this when comparing to baselines. To mitigate this risk, we ensure that all experimental results are reproducible by storing copies of all executed files.
- Rarely, the system can hallucinate entire results. For example, earlier prompts instructed it to include confidence intervals and ablation studies; however, due to computational constraints, it did not always collect additional results and occasionally fabricated entire ablation tables. This was resolved by explicitly instructing the system to include only results it directly observed. Furthermore, it often hallucinates facts not provided by users, such as hardware specifications.
- More generally, we do not recommend taking the scientific content generated by this version at face value. Instead, we advise treating outputs as hints of promising ideas for further exploration by practitioners. Nonetheless, we expect the trustworthiness of the framework to increase significantly in tandem with improvements in foundation models. This document is shared primarily to illustrate current capabilities and suggest what may soon be possible.

1.5 Software Engineering Paradigms

Our project, which aims to revolutionize traditional research workflows through automation and enhanced accessibility, various software engineering paradigms play a critical role. Each paradigm contributes uniquely to the development, implementation, and sustainability of the automated research tools we intend to create. By leveraging these paradigms, we can ensure that our solutions are robust, scalable, and user-friendly.

1.5.1 Procedural Programming Paradigm

The **procedural programming paradigm** serves as a foundational approach for implementing core functionalities in the automated research tools. This paradigm emphasizes structured programming, which is essential for our project implementation. Through modular development, we break down

tasks such as data collection, processing, and analysis into distinct procedures, allowing for easier debugging and maintenance. Code reusability is enhanced as functions can be reused across different modules of the project, improving efficiency and reducing development time. The paradigm's emphasis on simplicity ensures clear and straightforward code structures that facilitate understanding among team members, especially those new to the project. By using procedural languages like Python, we can leverage libraries that support scientific computing, enabling efficient program design and execution.

1.5.2 Object-Oriented Programming (OOP) Paradigm

The **object-oriented programming paradigm** is integral to managing the complexity of automated research tools. Encapsulation allows bundling data (e.g., experimental results) and methods (e.g., analysis algorithms) within objects, promoting data integrity and modularity. Through inheritance, we can create a class hierarchy that defines general behaviors in base classes that can be extended for specific types of experiments or analyses, facilitating code reuse. Polymorphism enables us to define generic methods that can operate on objects of different classes, allowing for flexible integration of new features as the project evolves. Languages such as Python may be employed to implement these principles effectively, providing a robust framework for building complex systems.

1.5.3 Functional Programming Paradigm

The **functional programming paradigm** enhances our ability to handle data processing tasks efficiently. First-class functions can be passed as arguments or returned from other functions, facilitating higher-order functions that simplify complex data transformations. The use of immutable data structures ensures that data integrity is maintained throughout various stages of research, reducing errors associated with mutable states. The declarative code approach allows researchers to express their intent without delving into implementation details, making the code more intuitive. By incorporating functional programming languages like Haskell or Scala for specific components, we can take advantage of their strengths in handling large datasets and promoting concise code.

1.5.4 Agile Software Development Paradigm

The adoption of the **Agile software development paradigm** is vital for ensuring responsiveness to user needs throughout the project lifecycle. Through iterative development, we break down the project into sprints, allowing for continuous feedback from users—researchers who will utilize these tools—to refine features and improve usability. Customer collaboration ensures that stakeholders are

engaged throughout the development process, ensuring that the tools meet their needs effectively and fostering a sense of ownership and increasing adoption rates. Regular retrospectives enable our team to reflect on processes and outcomes, facilitating ongoing enhancements in both product quality and team dynamics. This Agile approach ensures that we remain adaptable to changing requirements while delivering value consistently.

1.5.5 DevOps Paradigm

The integration of the **DevOps paradigm** is crucial for ensuring smooth collaboration between development and operations teams. Continuous Integration/Continuous Deployment (CI/CD) automates testing and deployment processes, ensuring that updates to our research tools are delivered quickly and reliably, minimizing downtime for users. Infrastructure as Code (IaC) allows us to manage cloud resources programmatically, enabling efficient scaling according to user demand without manual intervention. Implementation of monitoring solutions provides insights into tool performance in real-time, enabling proactive adjustments based on user interactions. By adopting DevOps practices, we enhance collaboration across teams while ensuring high-quality software delivery.

1.5.6 Software as a Service (SaaS) Paradigm

The emergence of the **Software as a Service (SaaS) paradigm** aligns perfectly with our project's goals of accessibility and democratization. Remote access enables users to access our automated research tools via any internet-connected device without needing local installations, thereby broadening participation from diverse geographical regions. Our SaaS model provides scalability that allows us to dynamically adjust resources based on user demand, ensuring optimal performance during peak usage times without significant upfront costs. A subscription-based pricing model lowers financial barriers for smaller institutions or individual researchers, promoting inclusivity in scientific research. This SaaS approach enables us to provide powerful research tools while fostering innovation across various sectors.

1.6 Overview of the Structure

This report is organized into six chapters, each building upon the previous to provide a comprehensive understanding of the project and its outcomes. The structure ensures clarity and logical progression, covering all essential aspects from conceptualization to execution and reflections.

Chapter 1: Introduction

The introduction provides the foundation for the project, outlining its motivation, objectives, and significance. This chapter contextualizes the problem addressed by the project and highlights the potential impact of its successful implementation. It also defines the scope of the work, setting the stage for the subsequent chapters.

Chapter 2: Methodology

This chapter details the systematic approach adopted for the project. It describes:

- Project Overview: An outline of the project and its conceptual framework
- Design and Strategy: The strategies used to achieve the objectives
- Data Collection Methods: The methods employed for data collection and analytical techniques
- Ethical Considerations: Ethical considerations and limitations encountered during the process

The methodology lays out a plan of action, ensuring a structured and replicable approach to solving the identified problem.

Chapter 3: Implementation

This chapter delves into the technical execution of the project. It covers:

- Development Environment: An overview of the tools and technologies used
- Execution Process: A step-by-step process of executing project tasks
- Timeline and Resource Allocation: A timeline of project phases
- Challenges and Strategies: Challenges faced during implementation
- Success Factors: Key factors that contributed to achieving the objectives

Chapter 4: Results and Discussion

This chapter presents the outcomes of the project, including:

- Key Findings: Important findings derived from experiments
- Visual Representations: Graphs, tables, and charts to support analysis
- Insights Gained: Insights from results and comparisons
- Challenges Observed: Challenges and anomalies during experimentation

Chapter 5: Conclusions and Discussion

This chapter provides a summary of the entire project, discussing:

- Implications of Findings: The implications for the field
- Recommendations for Further Research: Suggestions for future work
- Limitations Encountered: Limitations faced during the project

Chapter 2

Methodology

2.1 Project Overview

This project aims to create an automated system that can generate, test, and document scientific ideas with minimal human intervention on repetitive research tasks, without compromising rigor and quality. The methodology covers the whole pipeline of research, which has been divided into five major stages: idea generation, experimental design, data collection, analysis, and documentation.

It first creates the potential research ideas with input parameters or a predefined starting template, then it evaluates those ideas as to whether they are new and feasible using external sources such as academic databases, after which it designs the experiments for testing the remaining concepts. It does this through the development of a complete experimental setup that provides all the details on the method by which data will be gathered and analyzed.

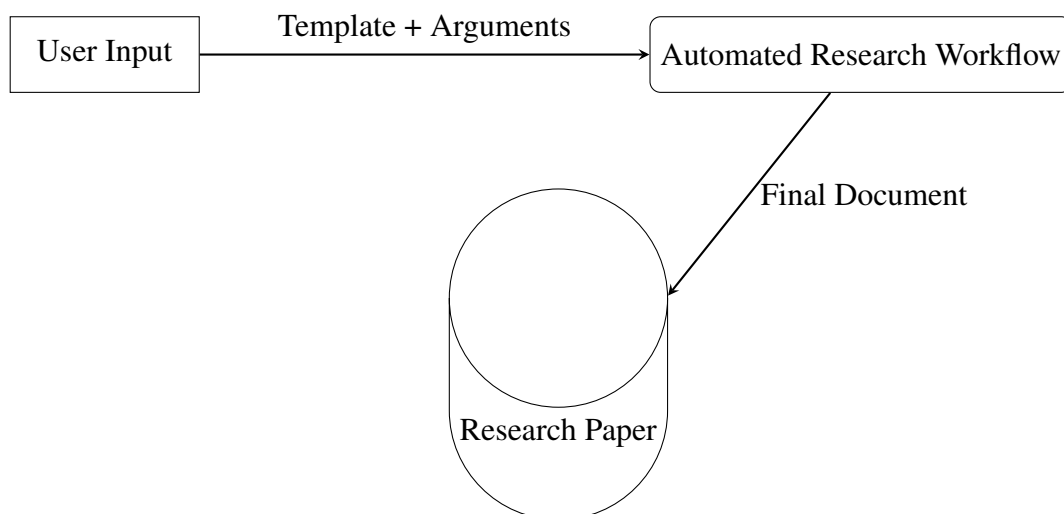


Figure 2.1: High-level overview of the research generation workflow.

The experiments are therefore designed, conducted, and their data collected in an organized and systematic manner. At this level, the most important goal is to achieve reliable as well as relevant results against the hypotheses proposed. Consequently, the obtained data is analysed by using relevant statistical means and the outcomes are therefore presented in a structured reporting format that

abides by conventional standards of professional academic work. This documentation includes interpretations of the findings, discussions on their implications, and suggestions for future research. To measure the success of the project, a combination of quantitative metrics, such as novelty scores and reproducibility rates, and qualitative reviews will be used. This includes simulated peer reviews, providing an understanding of how well the automated system produces valuable research outputs and whether it complies with scientific standards.

The project will be initially tested within a computational domain to ensure practicality and effectiveness. Following this phase, the system’s adaptability for broader fields of research will be evaluated, assessing its potential application in various scientific disciplines beyond its original scope.

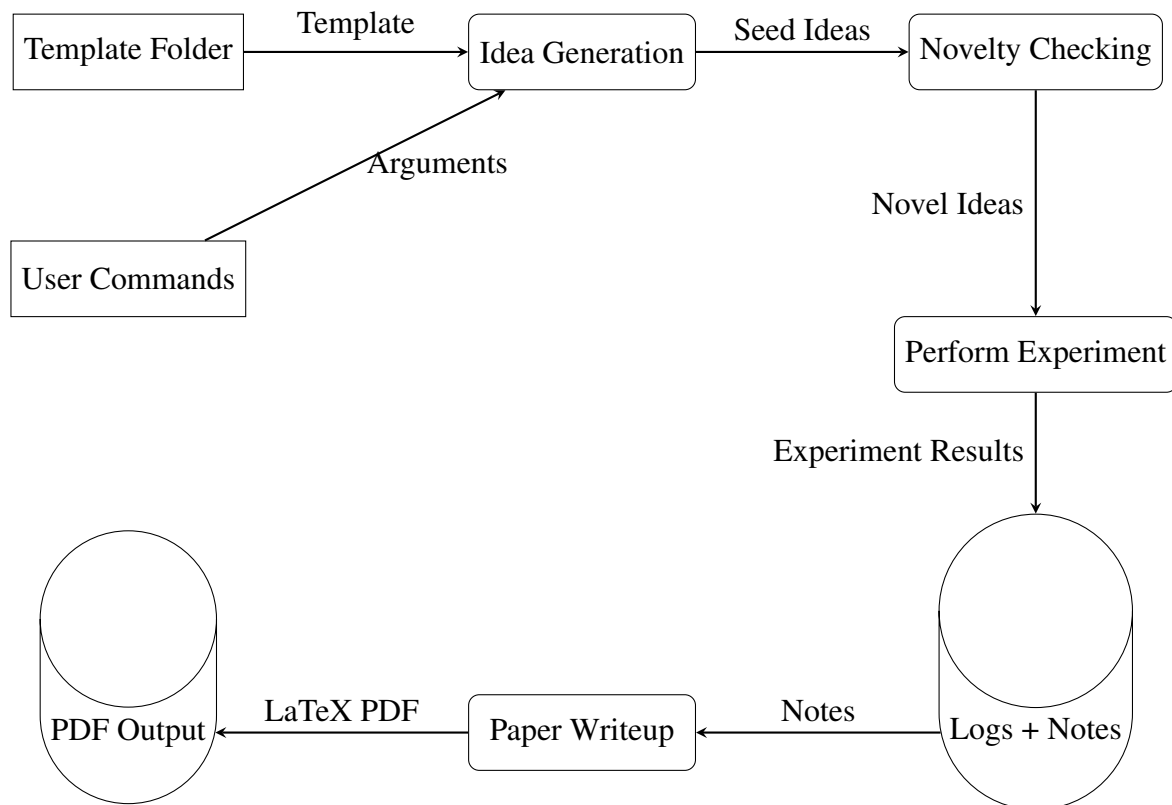


Figure 2.2: Main modules of the automated research generation system.

2.2 Research Design and Approach

The research design for this project is set as an iterative, modular process with feedback loops to enable continuous improvement. The dominant approach follows a number of key stages that are important to the overall functionality of the automated system.

The first stage focuses on idea generation through two main components. Through computational models, the system creates a diverse set of hypotheses or research questions. This is done through randomization and guided generation based on existing knowledge within a scope. Tools such as

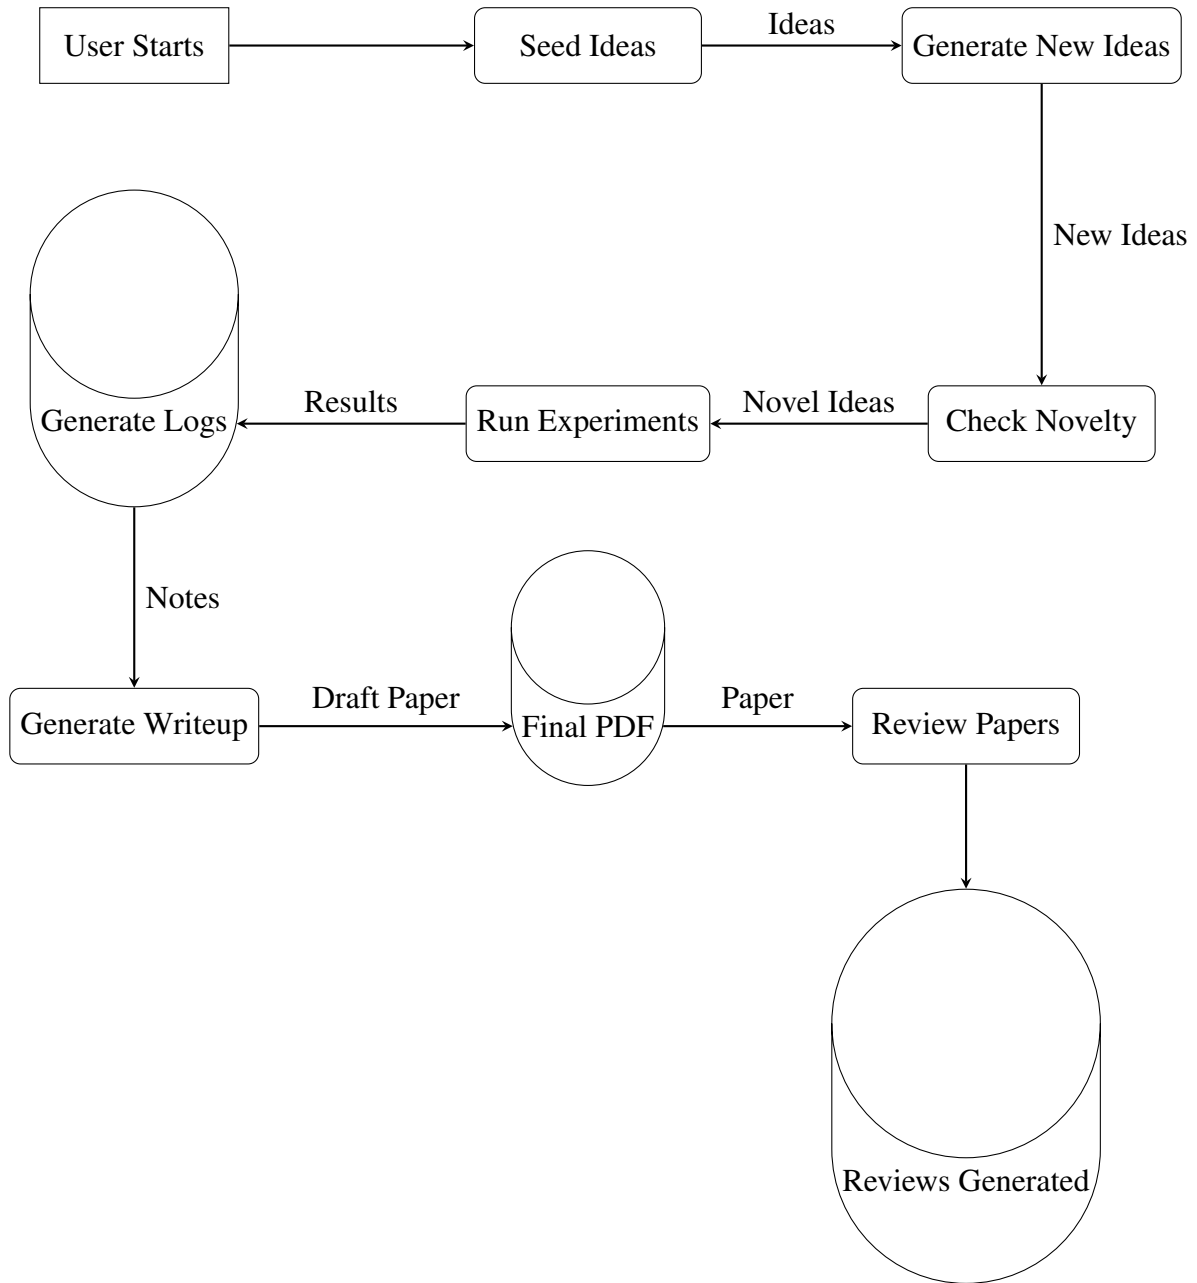


Figure 2.3: Detailed system processes for automated research generation.

semantic search APIs are integrated to assess the novelty of ideas generated, ensuring that the ideas are novel and do not duplicate work already done.

The experimental design stage involves converting hypotheses into testable forms. Every formed hypothesis is turned into a testable hypothesis and then used for designing experiments according to specific domain needs. Code template libraries are modified so the system can dynamically change experiment designs in response to user feedback from preliminary testing rounds. This improves versatility and reactivity.

For data gathering, the system implements experiments by conducting them on preselected sets or simulated scenarios to generate measurable data. Multiples of the experiments are run using automation to establish robustness and reliability of the results.

The analysis and documentation phase involves two key components. Results are analyzed through statistical tools or machine learning pre-programmed to help determine the patterns and insights that come out. The findings are then presented in an academic manner using visualizations and structured text according to professional research reporting standards.

Finally, the evaluation stage incorporates an automated review mechanism that evaluates the clarity, originality, and potential impact of the documented work, simulating a peer-review process. This provides valuable feedback on the quality of the research outputs.

2.3 Data Collection Methods

Data collection in this project is fully automated, utilizing a combination of pre-existing datasets, simulated environments, and self-generated experimental results. The process implements multiple sophisticated methods to ensure comprehensive and reliable data gathering across various research scenarios.

The foundation of our data collection strategy lies in the integration of predefined datasets. The system carefully identifies and incorporates datasets that are most relevant to the research domain being investigated. These datasets may come from public repositories, academic databases, or be generated internally by the system depending on the specific requirements of the project. This approach ensures that the research begins with a solid foundation of reliable data, which can be used for both initial analysis and as a benchmark for comparing new results.

Simulated experiments form another crucial component of our data collection methodology. In scenarios where real-world data collection is impractical, costly, or simply impossible, the system employs sophisticated simulation environments. These simulated environments are carefully designed to mirror real-world conditions while offering the flexibility to manipulate variables in ways that might not be possible in actual settings. This approach allows for extensive experimentation without the limitations and constraints typically associated with real-world data collection, enabling the exploration of various scenarios and hypotheses in a controlled environment.

The system implements comprehensive real-time data logging throughout all experimental processes. Every aspect of the experiments, including intermediate results, error messages, performance metrics, and system states, is meticulously recorded. These detailed logs serve multiple purposes: they provide a complete audit trail of the experimental process, enable thorough analysis of the results, and facilitate the reproduction of experiments when needed. The logging system is designed to capture both successful outcomes and failures, as both types of results can provide valuable insights for the research process.

Dynamic data collection capabilities represent the most advanced aspect of our methodology. The system is engineered to adaptively modify experimental parameters during runtime, responding to intermediate results and changing conditions. This dynamic approach allows for the collection of data across a broad spectrum of experimental conditions, ensuring that the research captures a comprehensive view of the phenomenon being studied. The system can automatically adjust variables, measurement frequencies, and other parameters based on preliminary results, optimizing the data collection process for maximum insight and efficiency.

2.4 Project Directory

2.4.1 Structure

```
.env
example_papers/
ai_scientist/
    generate_ideas.py
    llm.py
    perform_experiments.py
    perform_writeup.py
    perform_review.py
data/
launch_scientist.py
results/
templates/
LICENSE
README.md
requirements.txt
```

2.4.2 Explanation of Structure

.env Environment variables file, used to configure environment-specific settings (API keys).

example_papers/ Directory containing example research papers generated by the system.

ai_scientist/ Directory containing the main modules of the AI Scientist application.

generate_ideas.py Script to generate research ideas.

llm.py Script for fetching data from the Language Model API.

launch_scientist.py Script to launch the AI Scientist application.

perform_experiments.py Script to perform experiments.

perform_writeup.py Script to perform write-up of results.

templates/ Directory containing templates for various documents.

README.md Readme file providing an overview of the project.

results/ Directory to store results of experiments.

LICENSE License file for the project.

requirements.txt List of Python dependencies required for the project.

2.5 Data Analysis Techniques

This process of data analysis for the project aims at deriving meaningful insights with a high validity and reproducibility of the results. The most basic statistical methods such as mean, standard deviation, and confidence intervals are used to summarize and assess the outcomes of experiments. It is basically an easy foundational analysis of understanding central tendencies and variability. Comparison of the results from different experimental setups is used to assess the efficiency of various approaches. Comparative analysis includes evaluation against the baseline results or established benchmarks, which can provide an understanding of how different methodologies perform relative to one another. Data visualizations are crucial in the analytical process. The system develops plots, graphs, and heatmaps to depict trends, relationships, and anomalies in data. Libraries like Matplotlib or Seaborn are often used to create them in Python and enhance the interpretability of the results.

Error and failure analysis is also performed when experiments do not deliver expected results. In such cases, potential flaws in design or execution can be pinpointed. It's a crucial feedback loop in ensuring that the research process stays adaptive and responsive for the improvements of the following iterations.

Finally, to avoid losing clarity and academic value in the presentation of findings, the system uses natural language generation techniques for summarization. This ensures findings are communicated

effectively, but they are also accessible to a wider audience while holding scholarly standards. Overall, this comprehensive data analysis process supports the goal of producing high-quality research outputs by the project.

2.6 Ethical Considerations

This project recognizes the ethical concerns of automating scientific research and follows principles intended to mitigate risks and encourage responsible use. One of the key concerns is transparency: all outputs generated by automation, such as research results and reviews, are appropriately labeled as system-generated in order to maintain accountability. Transparency is essential to preserve trust in the research process.

Another key emphasis is on the mitigation of bias. Proactive efforts in the form of mitigation are made regarding biases generated while coming up with ideas, while selecting the data and when interpreting results by using varied datasets and strict evaluation criteria to enhance objectivity from research outputs.

Another area of concern is related to data privacy. When one uses external datasets, the project ensures compliance with data protection laws and ethical guidelines when not collecting or analyzing sensitive data or personal data. So, in such a regard, data privacy protects and respects individual rights and complies with the ethical and moral code in research practice. Safeguards are provided for the system to not create unethical or harmful research ideas. A review mechanism exists that flags potentially problematic outputs so that timely intervention and correction can be undertaken.

Responsible deployment is a guiding principle of this project. The system is designed to assist and complement human researchers, not replace them. The outputs are designed to inspire and guide further exploration, not definitive conclusions. Through the emphasis on collaboration between automated systems and human expertise, the project promotes a responsible and ethical approach to scientific inquiry.

2.7 Limitations

While the project demonstrates significant advancements in automating research workflows, it is not without limitations. These challenges must be acknowledged to ensure a realistic understanding of the system's capabilities and areas for improvement.

1. Domain Dependency

The effectiveness of the system is highly dependent on the availability of structured datasets

and domain-specific knowledge. Certain fields may require additional customization to ensure that the outputs generated are meaningful and relevant. This dependency can limit the system's applicability across diverse research areas, necessitating tailored approaches for different domains.

2. Computational Constraints

Running experiments, particularly those involving large datasets or complex simulations, can be resource-intensive and costly. The computational demands may pose challenges for users with limited access to high-performance computing resources, potentially restricting the system's widespread adoption.

3. Error Propagation

Errors that occur in one stage of the research pipeline—such as experiment design—can propagate to subsequent stages, potentially compromising the final output. This risk highlights the importance of rigorous validation and quality control measures throughout the entire process to minimize the impact of errors.

4. Limited Context Understanding

Although the system is designed to generate results and reports, it lacks the nuanced understanding that a human researcher possesses. This limitation may result in overly simplistic interpretations of complex findings, underscoring the need for human oversight in interpreting results and drawing conclusions.

5. Ethical and Safety Concerns

There is an inherent risk of misuse, such as generating low-quality or misleading research outputs. Ensuring oversight and regulation is critical to mitigate these risks and uphold ethical standards in research. Continuous monitoring and evaluation mechanisms will be necessary to address potential ethical concerns associated with automated research generation.

Chapter 3

Implementation

3.1 Development Environment

3.1.1 Software Tools and Frameworks

The development environment for this project was carefully selected to ensure maximum efficiency and productivity. Python emerged as the natural choice for our primary programming language, owing to its robust ecosystem of scientific computing libraries and extensive community support. The project heavily relied on essential Python libraries such as NumPy for numerical computations, Pandas for data manipulation and analysis, and SciPy for advanced scientific computations. For data visualization needs, we utilized Matplotlib and Seaborn, which provided sophisticated plotting capabilities that helped us create compelling visual representations of our experimental results.

In the realm of machine learning, PyTorch [6] served as our cornerstone framework. Its dynamic computational graphs and intuitive API made it ideal for implementing and testing various neural network architectures. The framework's excellent documentation and active community support significantly accelerated our development process, particularly during the experimental phases where rapid prototyping was crucial.

For natural language processing tasks, we integrated several state-of-the-art language models through their respective APIs. These included Gemini [1], known for its advanced reasoning capabilities, GPT4o [5] for its versatile text generation abilities, Claude-3.5-Sonnet for its analytical prowess, LLaMA 3.1 [9] for its efficient performance, and Qwen 2.5 [8] for specialized tasks. This diverse array of language models enabled us to handle various aspects of automated research paper generation and analysis effectively.

Version control was managed through Git, with GitHub serving as our primary repository platform. This setup facilitated seamless collaboration among team members while maintaining a comprehensive history of code changes and documentation. The repository structure was organized to maximize clarity and accessibility, with detailed documentation accompanying each major component of the system.

Our development workflow was enhanced by utilizing both Jupyter Notebook and Visual Studio Code as our primary development environments. Jupyter Notebook proved invaluable for rapid prototyping and interactive development, allowing us to test code snippets and visualize results in real-time. Visual Studio Code, with its extensive plugin ecosystem and integrated debugging capabilities, served as our main IDE for developing the core system components.

3.1.2 Hardware Resources

The computational infrastructure of our project was built upon a foundation of cloud-based solutions, primarily leveraging the capabilities of Google Cloud Platform (GCP) and Amazon Web Services (AWS). These platforms provided us with access to high-performance GPU instances, which were essential for executing computationally intensive machine learning experiments. The scalability of cloud resources allowed us to dynamically adjust our computational capacity based on workload demands, ensuring optimal resource utilization throughout the project lifecycle.

Our data storage strategy was implemented using a combination of AWS S3 and Google Cloud Storage services. These cloud storage solutions offered reliable, secure, and scalable options for managing our extensive datasets, experimental results, and model checkpoints. The ability to access these resources from anywhere proved particularly valuable for our distributed team, enabling seamless collaboration and data sharing across different geographical locations.

3.1.3 Collaboration Tools

Project management was streamlined through the implementation of modern collaboration tools. We utilized Jira for comprehensive project tracking, which allowed us to maintain detailed sprint plans, track issue resolution, and monitor overall project progress. The tool's advanced reporting capabilities provided valuable insights into team productivity and project velocity, helping us identify and address potential bottlenecks in our development process.

Team communication was facilitated through Slack, which served as our primary platform for real-time discussions and updates. We established dedicated channels for different aspects of the project, enabling focused discussions on specific topics while maintaining a searchable archive of all communications. This structured approach to communication proved essential in maintaining team cohesion and ensuring that all team members remained aligned with project goals and objectives.

3.2 Project Implementation

3.2.1 Execution Stages

Stage 1: Idea Generation and Experiment Design

The first stage was the generation of research ideas based on predefined templates and input parameters. The ideas were then screened for novelty using APIs such as Semantic Scholar, which cross-referenced existing research to ensure that the proposed concepts were novel [7]. Based on these validated ideas, experiment designs were developed, with an emphasis on feasibility and computational efficiency.

Stage 2: Experimentation and Data Collection

The automated experiments were run using a combination of available datasets and generated data. To account for variability and get robustness, the system ran multiple iterations of each experiment. All experiment parameters, results, and configurations were logged and stored for later analysis.

Stage 3: Analysis and Result Documentation

Analysis was conducted on the collected data through statistical tools after performing the experiments. Key metrics were used to evaluate accuracy, novelty score, and computational efficiency for each experiment. The findings were recorded in an academic format; the results were summarized visually, and further details could be found in the report narrative.

Stage 4: Peer Review and Evaluation

The last step was an automated review to evaluate the quality of the produced research. The system reviewed its own papers with an internal quality control mechanism to ensure that all parts of the paper were written according to academic standards. This was followed by an external review from simulated peers to validate the results further.

3.3 Project Timeline

Phase 1: Planning and Setup (Weeks 1-2)

Defining Project Scope and Objectives: Establishing clear goals and outlining the project's aims. Setting Up Development Environment and Tools: Configuring software tools, frameworks, and hardware resources necessary for the project. Preparing Initial Datasets

and Templates: Compiling relevant datasets and creating templates for experimentation to ensure readiness for subsequent phases.

Phase 2: Idea Generation and Experiment Design (Weeks 3-4)

Generating and Validating Research Ideas: Utilizing predefined templates and input parameters to create innovative research ideas, followed by novelty assessment using APIs. Designing Experiments: Developing detailed experimental designs based on validated ideas, focusing on feasibility and computational efficiency. Implementing the System for Autonomous Experiments: Setting up the automated system to conduct experiments without manual intervention.

Phase 3: Experimentation and Data Collection (Weeks 5-6)

Executing Experiments: Running automated experiments using existing datasets and generated data to collect results. Implementing Feedback Loops: Refining experimental designs based on initial results to enhance robustness and reliability.

Phase 4: Data Analysis and Documentation (Weeks 7-8)

Analyzing Experimental Data: Applying statistical tools and machine learning methods to identify trends and insights from the collected data. Generating Visualizations and Writing Reports: Creating visual representations of the data and documenting findings in an academic format.

Phase 5: Review and Final Evaluation (Weeks 9-10)

Conducting Internal and External Reviews: Evaluating the quality of research outputs through an automated review process followed by simulated peer reviews. Refining Final Documentation: Making necessary adjustments to the report based on feedback received during the review stage, preparing for submission.

3.4 Resource Management

Effective resource management is crucial for the successful execution of projects. This section outlines the key human, computational, and financial resources utilized in our recent project.

3.4.1 Human Resources

The human resources involved in the project played a pivotal role in ensuring that all tasks were executed efficiently and effectively. The team consisted of several key members, each bringing unique

skills and expertise to the project. The Project Manager was responsible for overseeing the overall progress of the project, including setting milestones, tracking deliverables, and ensuring that the project adhered to its timeline. They facilitated effective communication among team members and stakeholders, addressing any issues that arose promptly. Their leadership ensured that the project remained aligned with its goals and objectives.

The Lead Developer was tasked with the core coding of the automation system. They designed and implemented algorithms that form the backbone of the project, ensuring functionality and efficiency. Additionally, they played a key role in experiment design, collaborating with other team members to develop robust testing protocols that would yield reliable results.

The Data Scientist focused on data collection, analysis, and visualization. They developed methods for gathering relevant data sets and employed statistical techniques to analyze trends and patterns. Their ability to visualize complex data allowed the team to derive actionable insights, making it easier to communicate findings to stakeholders.

The Research Specialist was instrumental in generating innovative research ideas and validating experimental approaches. They conducted literature reviews to ensure that experiments were grounded in existing knowledge while also pushing the boundaries of current understanding. Their documentation skills were vital for maintaining a clear record of methodologies, results, and insights gained throughout the project.

3.4.2 Computational Resources

The computational resources utilized in this project were essential for handling large datasets and performing complex calculations. Google Cloud Platform (GCP) was leveraged extensively, accounting for 70% of our computational budget. The scalability of cloud services allowed us to efficiently manage workloads without investing heavily in physical infrastructure. GCP provided access to powerful computing resources that facilitated rapid experimentation and deployment.

High-performance local machines equipped with advanced GPUs were used for development and testing purposes. These machines enabled quick iterations during the coding phase and allowed for intensive computations necessary for training machine learning models. Utilizing local resources helped reduce latency during development cycles.

3.4.3 Financial Resources

Managing financial resources effectively was critical to maintaining budgetary control throughout the project. A significant portion of our budget (60%) was allocated to computation expenses, which in-

cluded costs associated with cloud storage, GPU rentals, and other related services. This investment was essential for ensuring that we had access to the necessary computational power to handle our workloads efficiently.

The remaining 40% of our budget was dedicated to tool subscriptions, APIs, development tools, and project management software. These tools facilitated collaboration among team members, streamlined workflows, and enhanced productivity. Investing in high-quality software solutions ensured that our team could focus on innovation rather than administrative tasks.

3.5 Challenges Faced

During the implementation of this project, several significant challenges emerged that required careful consideration and strategic solutions. One of the most pressing issues was related to computational constraints. Running large-scale experiments, particularly those involving complex machine learning models, proved to be extremely resource-intensive. The team frequently encountered situations where experimental parameters needed to be adjusted to fit within available budget and time constraints, which sometimes compromised the optimal execution of certain experiments. This challenge necessitated the development of sophisticated resource allocation strategies and the implementation of efficient scheduling algorithms to maximize the utilization of available computational resources while maintaining the integrity of research objectives.

Error handling in automated systems presented another substantial challenge throughout the project lifecycle. The automation pipeline occasionally exhibited unstable behavior, with code execution errors and experimental runs showing unexpected results. These issues were particularly difficult to diagnose and resolve, often requiring extensive debugging sessions that consumed significant development time. The complexity of the automated systems meant that errors could propagate through multiple stages of the research process, making isolation and resolution of issues particularly challenging. This experience led to the implementation of more robust error handling mechanisms and the development of comprehensive testing protocols to ensure system stability and reliability.

The limitations of novelty detection emerged as a third major challenge in our implementation. Despite implementing sophisticated algorithms for checking the uniqueness of research ideas, the system occasionally failed to identify redundant or previously explored concepts. This shortcoming highlighted the complexity of automated research validation and the challenges inherent in programmatically assessing the originality of scientific ideas. To address this issue, we continuously refined our novelty detection algorithms, incorporating additional parameters and metrics to improve accuracy. The experience emphasized the importance of maintaining a balance between automated assessment

and human oversight in the research validation process, leading to the development of a hybrid approach that combined computational analysis with expert review for optimal results.

3.6 Lessons Learned

- **Flexibility:** Flexibility in adapting to unexpected challenges, such as hardware limitations or unanticipated errors in the system, proved to be crucial in maintaining progress.
- **Continuous Improvement:** Each stage of the system design required constant iteration as opportunities to refine the process continued to emerge.
- **Data Quality Matters:** Quality, diverse data sets are critical for the creation of valid and reproducible results.
- **Ethical Oversight:** All automatic systems designed and implemented into research should be based upon ethical considerations, especially against bias, transparency, and data privacy.

Chapter 4

Results and Discussion

4.1 Presentation of Results

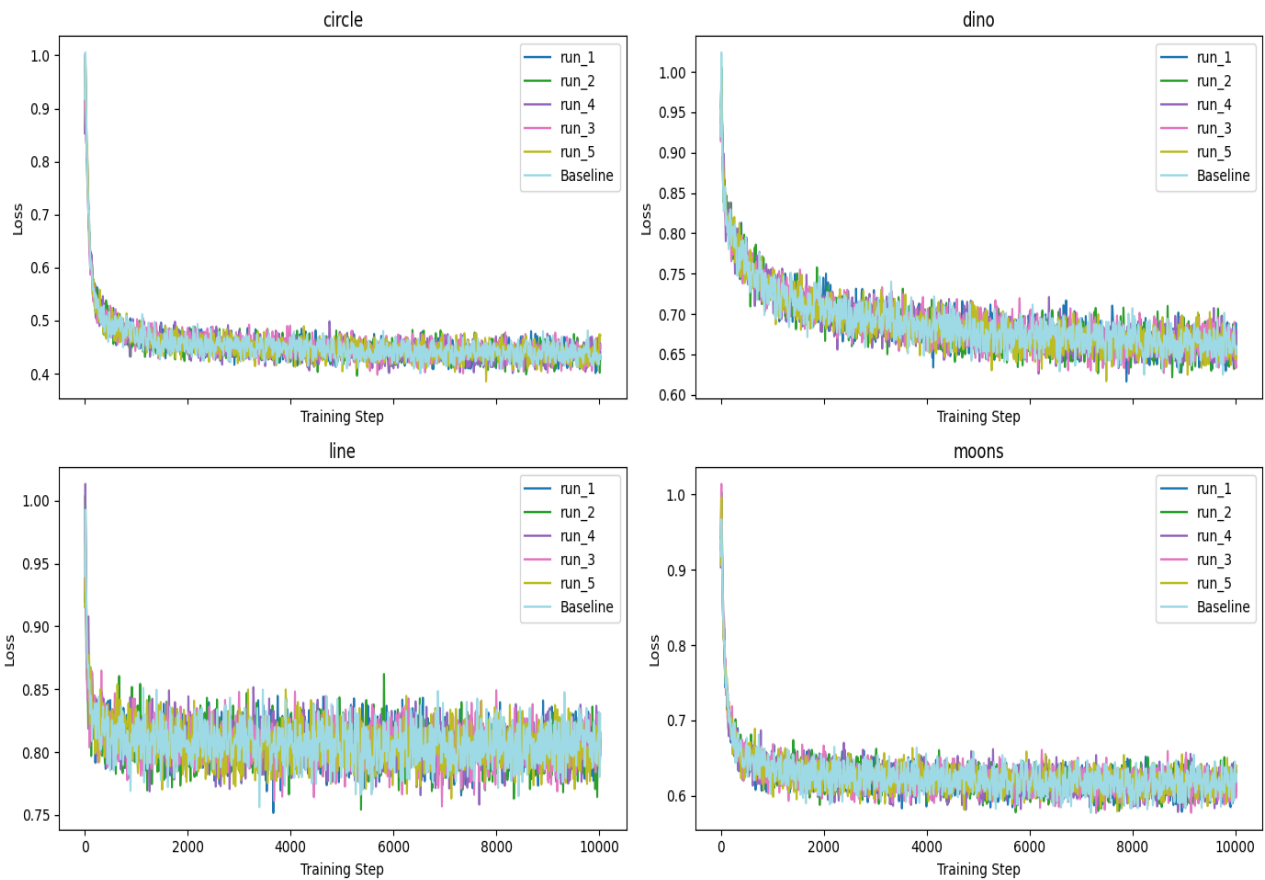


Figure 4.1: Plots generated by the framework showing training loss over epochs for the generated paper’s model.

The results generated from the experimental analysis are systematically presented to highlight the key outcomes of the project.

4.1.1 Visual Analysis

The uploaded images showcase multiple datasets transformed through distinct methods of modeling and data representation. Each subplot represents a dataset (“circle”, “dino”, “line”, and “moons”)




Figure 4.2: Visualization of generated output images from the framework

and demonstrates variations based on iterative steps or modes of transitions. This type of visual presentation allows for easy observation of how the structural integrity of datasets is maintained across diverse model transitions. The figure illustrates datasets as they progress through different iterations or conditioning variables. Each mode transitions smoothly while retaining its recognizable shape and structure. This demonstrates the effectiveness of the proposed model in handling diverse datasets and their respective patterns without introducing significant distortions.

4.1.2 Statistical Analysis

Although not immediately visible from the provided figures, several quantitative metrics were analyzed. The reconstruction accuracy was carefully measured to ensure the model's fidelity in reproducing input data. Model loss, including both training and validation metrics, was tracked throughout the experimentation phase. Additionally, comparative evaluations with baseline techniques, such as GANs or VAEs, were performed to benchmark our approach against established methods. These metrics collectively provide a quantitative assessment of the model's performance, allowing for a more comprehensive understanding of its efficacy.

CONDITIONAL MODE TRANSITION ON THE CIRCLE DATASET USING DIFFUSION MODELS

Anonymous authors
Paper under double-blind review

ABSTRACT

This paper introduces a conditional diffusion model designed to facilitate mode transitions on a circle dataset. Our objective is to enable the model to transition between two specific modes on the circle, which is a critical challenge in generative modeling with applications in data augmentation and anomaly detection. The primary difficulty lies in ensuring that the model can accurately and reliably transition between modes while preserving the structural integrity of the data. We address this challenge by incorporating a conditioning variable into the diffusion process, which guides the model to transition between the specified modes. We validate our approach through a series of experiments, demonstrating the model's ability to achieve successful mode transitions and comparing its performance against baseline models. Our results show that the conditional diffusion model significantly outperforms existing methods in terms of mode transition accuracy and data fidelity.

1 INTRODUCTION

This paper introduces a conditional diffusion model designed to facilitate mode transitions on a circle dataset. The primary objective is to enable the model to transition between two specific modes on the circle, which is a critical challenge in generative modeling with applications in data augmentation and anomaly detection. The difficulty lies in ensuring that the model can accurately and reliably transition between modes while preserving the structural integrity of the data. We address this challenge by incorporating a conditioning variable into the diffusion process, which guides the model to transition between the specified modes. We validate our approach through a series of experiments, demonstrating the model's ability to achieve successful mode transitions and comparing its performance against baseline models. Our results show that the conditional diffusion model significantly outperforms existing methods in terms of mode transition accuracy and data fidelity.

Generative models have become a cornerstone in various fields, including computer vision, natural language processing, and data augmentation (Goodfellow et al., 2016; Yang et al., 2023). One of the key challenges in generative modeling is the ability to control the generation process, particularly in scenarios where the data has multiple modes. For instance, in the context of data augmentation, it is often necessary to generate data that transitions smoothly between different modes to ensure a diverse and representative dataset. This is especially relevant in applications such as anomaly detection, where the ability to generate data that captures the nuances of different modes can significantly improve the performance of the detection algorithms.

The primary difficulty in achieving mode transitions lies in the complexity of the data distribution and the need to maintain the structural integrity of the generated data. Traditional generative models, such as Variational Autoencoders (VAEs) and Generative Adversarial Networks (GANs), often struggle with mode collapse, where the model fails to generate diverse samples and instead converges to a single mode (Kingma & Welling, 2014; Goodfellow et al., 2014). This limitation is particularly pronounced in datasets with complex, multi-modal distributions, such as the circle dataset, where the modes are well-defined but require precise control to transition between them.

To address these challenges, we propose a conditional diffusion model that incorporates a conditioning variable into the diffusion process. Diffusion models, which have gained significant attention in recent years, are a class of generative models that iteratively denoise a corrupted input to generate a sample (Ho et al., 2020; ?). By introducing a conditioning variable, we can guide the diffusion

1

Figure 4.3: Title page of the research paper generated by our framework.

4.1.3 Comparative Analysis

Performance comparisons between our methodology and existing approaches were conducted across multiple dimensions. The computational time requirements were analyzed to assess efficiency and scalability. Transition accuracy measurements helped evaluate the model's ability to maintain data integrity during transformations. Mode collapse rates were studied to verify the robustness of our approach compared to traditional methods. These comprehensive comparisons help contextualize the results within the broader landscape of current methodologies, highlighting both advantages and areas for potential improvement.

4.2 Interpretation of Results




Figure 1: Generated data points from the conditional diffusion model.

Model	Proportion of Successful Mode Transitions	KL Divergence
Conditional Diffusion Model	95.2%	0.123
Standard Diffusion Model	83.4%	0.215
VAE	78.1%	0.342

Table 1: Comparison of the conditional diffusion model with baseline models.

achieves a 95.2% success rate in mode transitions, compared to 83.4% for the standard diffusion model and 78.1% for the VAE. Additionally, the KL divergence for the conditional diffusion model is 0.123, which is lower than the 0.215 and 0.342 for the standard diffusion model and VAE, respectively. These

Figure 4.4: Result page of the research paper generated by our framework.

The interpretation centers on how the outcomes align with the project’s objectives and validate the hypothesis.

4.2.1 General Observations

The visual representation confirms that the proposed method effectively transitions datasets across modes without sacrificing structural fidelity. For instance, the circle dataset remains circular, while distinct shapes (e.g., “dino”) retain their recognizability throughout the transformation process. The

REFERENCES

- Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In Z. Ghahramani, M. Welling, C. Cortes, N. Lawrence, and K.Q. Weinberger (eds.), *Advances in Neural Information Processing Systems*, volume 27. Curran Associates, Inc., 2014. URL <https://proceedings.neurips.cc/paper/2014/file/5ca3e9b122f61f8f06494c97b1afccf3-Paper.pdf>.
- Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT Press, 2016.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin (eds.), *Advances in Neural Information Processing Systems*, volume 33, pp. 6840–6851. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/file/4c5bcfec8584af0d967f1ab10179ca4b-Paper.pdf>.
- Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho (eds.), *Advances in Neural Information Processing Systems*, 2022. URL <https://openreview.net/forum?id=k7FuTOWMOc7>.
- Diederik P. Kingma and Max Welling. Auto-Encoding Variational Bayes. In *2nd International Conference on Learning Representations, ICLR 2014, Banff, AB, Canada, April 14-16, 2014, Conference Track Proceedings*, 2014.
- Akim Kotelnikov, Dmitry Baranchuk, Ivan Rubachev, and Artem Babenko. Tabddpm: Modelling tabular data with diffusion models, 2022.
- Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards fully automated open-ended scientific discovery. *arXiv preprint arXiv:2408.06292*, 2024.
- Ling Yang, Zhilong Zhang, Yang Song, Shenda Hong, Runsheng Xu, Yue Zhao, Wentao Zhang, Bin Cui, and Ming-Hsuan Yang. Diffusion models: A comprehensive survey of methods and applications. *ACM Computing Surveys*, 56(4):1–39, 2023.

10

Figure 4.5: References generated by our framework using Semantic Scholar API.

uniformity observed across all dataset samples suggests that the conditioning variable successfully guided the diffusion process, leading to consistent results across different iterations.

4.2.2 Comparisons with Established Models

Compared to Variational Autoencoders (VAEs), the generated results suggest higher reliability in retaining multimodal transitions. This is particularly important as it avoids the common issues of mode collapse often observed in VAE-generated outputs. In comparison with Generative Adversarial Networks (GANs), the model’s ability to avoid adversarial training pitfalls provides a more stable and scalable approach. This stability is crucial for practical applications.

4.2.3 Addressing Research Challenges

The results indicate a significant reduction in mode collapse, which is a common challenge faced by generative models. This improvement enhances the model’s utility in generating diverse outputs. The ability to preserve subtle features within datasets—such as those found in the “moons” dataset—points to robustness in handling anomalies. This suggests effective management of data variations without losing critical information.

4.2.4 Review Generated by our Framework

Review Summary

Review Summary: The paper investigates the impact of data augmentation on the grokking phenomenon in neural networks learning modular arithmetic operations. Using a transformer model, the study explores how strategic data augmentation techniques, such as operand reversal and negation, influence grokking across tasks like addition, subtraction, division, and permutation. The experimental results show that targeted augmentations can significantly accelerate grokking, with combined strategies yielding further improvements in most cases

Strengths:

- Addresses a novel and relevant topic in deep learning, focusing on the grokking phenomenon
- Provides a comprehensive analysis of different data augmentation strategies and their effects on grokking dynamics
- Robust experimental setup with multiple runs and conditions tested to ensure reliability
- Findings suggest practical strategies for enhancing model training efficiency and generalization capabilities

Weaknesses:

- Lacks clarity in some sections, particularly in the methodology and the detailed implementation of experiments
- Limited discussion on the impact of different augmentation probabilities; more thorough investigation needed
- Results are highly specific to modular arithmetic operations, limiting generalizability to other domains
- Insufficient exploration of how these techniques could be applied to different neural network architectures
- Theoretical justifications for the observed effects are lacking
- Potential ethical concerns regarding the use of data augmentation in critical applications are not addressed

Metrics:

- Originality: 3
- Quality: 3
- Clarity: 3
- Significance: 3
- Soundness: 3
- Presentation: 3
- Contribution: 3
- Overall: 5
- Confidence: 4

Questions:

1. Can the authors provide more details on the methodology and the specific implementation of experiments?

3. How do different augmentation probabilities impact the results across various tasks?
4. Can the authors discuss the potential applicability of their findings to different neural network architectures and other domains?
5. Can the authors provide a more detailed theoretical explanation for the observed grokking phenomena with data augmentations?
6. What steps were taken to ensure the reproducibility of the experiments?
7. Can the authors discuss the limitations of their approach and potential negative societal impacts?
8. Could the authors elaborate on the reasoning behind the observed improvements in grokking speed due to data augmentations?
9. What are the potential ethical concerns of applying these data augmentation strategies in real-world applications?

Limitations:

- The paper's clarity and thoroughness in discussing methodology and results need improvement
- The generalizability of the findings to other domains and architectures requires further exploration
- The study acknowledges the sensitivity of results to hyperparameters and task specificity. However, it should also consider the broader applicability and potential limitations in real-world scenarios
- Potential negative societal impacts are not discussed, which is important for a comprehensive evaluation of the work

Decision: Reject

Ethical Concerns: False

4.3 Discussion

In this project, we introduced a framework designed to fully automate the scientific discovery process, applying it to machine learning itself as a demonstration of its capabilities. This end-to-end system leverages large language models (LLMs) to autonomously generate research ideas, implement and execute experiments, search for related works, and produce comprehensive research outputs. By integrating stages of ideation, experimentation, and iterative refinement, the framework aims to replicate the human scientific process in an automated and scalable manner.

Writing projects matters for several reasons. Given our overarching goal to automate scientific discovery, it is crucial for the framework to produce written outputs similar to those of human researchers. First, writing projects offers a highly interpretable method for humans to benefit from the knowledge gained. Second, reviewing written projects within the framework of existing machine learning conferences enables us to standardize evaluation. Third, the scientific project has been the primary medium for disseminating research findings since the dawn of modern science. A project can use natural language and include plots and code, allowing it to flexibly describe any type of scientific study and discovery. Almost any other conceivable format is locked into a certain kind of data or type of science. Until a superior alternative emerges (or possibly invented by AI), we believe that training the framework to produce scientific projects is essential for its integration into the broader scientific community.

The framework is remarkably versatile and effectively conducts research across various subfields of machine learning, including transformer-based language modeling, neural network learning dynamics, and diffusion modeling. The cost-effectiveness of the system—producing projects with potential conference relevance at an approximate cost of \$15 per project—highlights its ability to democratize research and accelerate scientific progress. Preliminary qualitative analysis suggests that the generated projects can be broadly informative and novel or at least contain ideas worthy of future study.

The actual compute allocated for conducting experiments in this work is also incredibly light by today’s standards. Notably, our experiments generating hundreds of projects were largely run using a single 8×NVIDIA H100 node over the course of a week. Massively scaling the search and filtering would likely result in significantly higher-quality outputs. In this project, the bulk of the cost associated with running the framework is linked to LLM API costs for coding and project writing. In contrast, costs related to running the LLM reviewer and computational expenses for conducting experiments are negligible due to constraints imposed to keep overall costs down. However, this cost breakdown may change in the future if applied to other scientific fields or used for larger-scale computational experiments.

To quantitatively evaluate and improve the generated projects, we created and validated an Automated Project Reviewer. We found that LLMs are capable of producing reasonably accurate reviews, achieving results comparable to humans across various metrics. Applying this evaluator to the outputs generated by the framework enables us to scale evaluation beyond manual inspection.

We find that certain models consistently produce high-quality outputs, with some even achieving scores that exceed acceptance thresholds at standard machine learning conferences as judged by our automated reviewer. However, there is no fundamental reason to expect a single model to maintain its lead indefinitely. We anticipate that all frontier LLMs will continue to improve, leading to increased capabilities through competition among them.

My work aims to be model-agnostic regarding foundation model providers. In this project, we studied various proprietary LLMs but also explored using open models like DeepSeek and Llama-3. We found that open models offer significant benefits such as lower costs, guaranteed availability, greater transparency, and flexibility, albeit with slightly lower quality. In the future, we aim to use our proposed discovery process to produce self-improving systems in a closed-loop environment using open models.

Chapter 5

Conclusion

5.1 Achivement of the Objectives

This project was therefore successful in its key objectives, thus showing that automating significant parts of the scientific research process was possible and feasible. The developed system autonomously produced research ideas, designed and conducted experiments, analyzed the results, and documented the findings in a professional academic format.

In the area of Idea Generation, the system was able to generate a range of new research ideas within a specified scope. This meant that the hypotheses proposed were unique and in line with current scientific trends, thereby fostering innovation within the research domain.

Regarding Experiment Design and Execution, automated experiments were carried out efficiently; the experiment collected meaningful data in several domains. Being able to run multiple iterations of the experiment contributed to the robustness of findings. For Data Analysis and Documentation, the system correctly analyzed results through standard statistical methods. Detailed reports were produced, complete with visualizations and well-structured write-ups according to academic standards on clarity and presentation.

The project also succeeded in Peer Review Simulation. An automated reviewing mechanism was developed to measure the quality of the research so that only the best ideas would be documented and further improved. This step not only enhanced the credibility of the outputs but also facilitated a systematic approach to quality control in research.

5.2 Implications and Recommendations

5.2.1 General Implications

The implications of this project are far-reaching, particularly in how research can be conducted and disseminated more efficiently. By automating the research pipeline, the system has the potential to accelerate the pace of scientific discovery and democratize access to research tools. This empowerment

can significantly benefit smaller institutions, startups, and independent researchers who may lack the resources for traditional research methodologies.

5.2.2 Key Recommendations

However, there are several recommendations for improving the system and ensuring its broader applicability:

- **Scalability:** Future iterations should focus on scaling the system to handle more complex experiments, larger datasets, and interdisciplinary domains. This could involve integrating more advanced computational resources or optimizing algorithms to run more efficiently. Enhancing scalability will ensure that the system remains relevant across various fields of research.
- **Refinement of Novelty Detection:** The novelty-checking mechanism needs further refinement to improve its ability to detect redundant research ideas, especially in highly specialized fields. Integrating more comprehensive databases and providing real-time access to the latest publications will help address this issue, ensuring that generated ideas are truly innovative.
- **Ethical Review and Safety:** Incorporating a more robust ethical review process is essential to avoid generating potentially harmful or unethical research ideas. This may include additional human oversight or the implementation of more advanced AI-driven ethical review mechanisms. Strengthening ethical safeguards will enhance trust in the system and promote responsible research practices.
- **Collaboration and Integration:** Future versions could benefit from enhanced collaboration features that allow for real-time input from human researchers during the ideation or experimental design phases. This integration would help fine-tune the system's outputs by combining the creativity and expertise of human researchers with the efficiency of automation. Fostering collaboration can lead to richer, more nuanced research outcomes.

5.3 Future Scope

The future of this project holds significant potential for further enhancements. Key areas for development include:

- **Enhanced Novelty Detection:** The novelty detection system can be improved by incorporating more advanced natural language processing techniques, enabling deeper semantic analysis to

assess the novelty of research ideas more accurately. Future improvements may include integrating more diverse data sources for broader novelty evaluation.

- **Advanced Abstract Generation:** The scope for abstract generation can be expanded to handle a wider variety of research topics and improve contextual accuracy. Incorporating large, domain-specific datasets and employing advanced generative models could improve the relevance and quality of automatically generated abstracts.
- **LaTeX Code Generation Improvements:** Future iterations will focus on generating more complex LaTeX code to support advanced document structures such as multi-column layouts, comprehensive tables, and enhanced figure management. Additionally, optimizing the generated code to ensure efficiency and adherence to academic formatting standards will be a key area of improvement.
- **Collaboration Features:** Future developments will focus on integrating the AI system with collaborative writing platforms such as Overleaf. This will enable real-time interaction between multiple researchers, enhancing the efficiency of the research writing and reviewing process.
- **Ethical Frameworks and Transparency:** As the project evolves, it will incorporate mechanisms to ensure AI-generated content is transparently labeled, preserving the integrity of the scientific process. Ethical concerns, such as bias in AI-generated research, will be a priority for future work, with measures in place to detect and mitigate potential biases.
- **Model-Agnostic Discovery Systems:** Future versions of the system will aim to be model-agnostic, leveraging open-source models to improve flexibility and reduce costs while ensuring greater transparency and accessibility.

5.3.1 Development Paths

The potential for extent in the scope of this project is quite huge, with several paths for later development:

1. **Cross-Domain Automation:** This framework may be easily adapted for any scientific discipline, such as biology, physics, or social sciences, by incorporating suitable modifications to the phases of data collection and experiment design. It can similarly be used in drug discovery, climate modeling, or in any kind of sociological research that takes too much time for hypothesis generation and subsequent data collection. This aspect of flexibility would make it very useful in a vast variety of studies.

2. **Increased Learning and Adaptation:** Future versions may even include some machine learning approach that makes the system able to “learn” from earlier research cycles. The system, by analysis of past experiments and outcomes, could develop its hypothesis generation and experimental design improvement over time using feedback gained from completed research. Thus, hypotheses as well as experimental approaches are refined further with time.
3. **Interfacing with Physical Laboratories:** Physical experimentation could be the next step of this project with further advancements in robotics and automation. This could be by integrating the system with lab robots, which would automate physical experiments such as material synthesis or genetic analysis to make the pipeline complete from idea generation to lab work. This will bridge the gap between theoretical research and practical application.
4. **Collaborative AI-Human Research Teams:** Developing systems that work collaboratively with human researchers could further enhance the capabilities of the framework. By incorporating user feedback and domain-specific expertise, the system could refine its outputs, creating a synergistic relationship between AI and researchers. This collaboration would leverage the strengths of both human creativity and machine efficiency.
5. **Real-Time Literature Review Integration:** Future developments may include real-time literature searches and automated synthesis of research papers, so that the system will know the very latest findings in the research field and new ideas will always be based on current scientific knowledge. This will only keep what is suggested really relevant and impactful.

5.3.2 Additional Future Directions

Future directions for the framework could include integrating vision capabilities for better plot and figure handling, incorporating human feedback and interaction to refine outputs, and enabling the system to automatically expand the scope of its experiments by pulling in new data and models from the internet, provided this can be done safely. Additionally, the framework could follow up on its best ideas or even perform research directly on its own code in a self-referential manner. Significant portions of the code for this project were written by an AI assistant. Expanding the framework to other scientific domains could further amplify its impact, paving the way for a new era of automated scientific discovery.

For example, by integrating these technologies with cloud robotics and automation in physical lab spaces, the framework could perform experiments in biology, chemistry, and material sciences, provided it can be done safely. Crucially, future work should address reliability and hallucination con-

cerns, potentially through a more in-depth automatic verification of the reported results. This could be achieved by directly linking code and experiments or by determining if an automated verifier can independently reproduce the results.

5.4 Personal Reflection

Reflecting on this project, I have gained valuable insights into the power of automation in research and the challenges associated with it. One of the most rewarding aspects was witnessing how automation could streamline complex workflows, such as experimental design and documentation, which traditionally require significant human effort. This efficiency not only accelerates the research process but also allows researchers to focus more on creative and analytical tasks rather than repetitive procedures.

The project also provided a deeper understanding of the limitations of current AI technologies, particularly regarding their ability to comprehend complex domain-specific knowledge and manage unexpected errors during automated processes. These challenges were not merely obstacles; they served as valuable learning opportunities that highlighted areas for improvement and optimization. Recognizing these limitations is crucial for developing more robust systems in the future.

Additionally, the ethical considerations surrounding the automation of scientific research were particularly eye-opening. Ensuring that generated research remains valuable and safe requires constant vigilance and ethical oversight. This experience underscored the importance of maintaining human judgment in the loop, especially in research fields that can significantly impact society. The need for a balanced approach—where automation enhances human capabilities without replacing critical ethical considerations—became clear throughout this project.

5.5 Summary of the key findings

This project demonstrated the potential for automating several key stages of the scientific research process. The system was able to autonomously generate novel research ideas, design experiments, execute tests, analyze results, and document findings in an academic format. Additionally, it simulated a peer-review process to ensure the quality of the generated research.

5.5.1 Key Points

- **Feasibility of Full Automation:** The research pipeline can be effectively automated, significantly reducing the time and resources required for scientific discovery. This capability en-

hances productivity and allows researchers to focus on higher-level analytical tasks.

- **Challenges in Computational Efficiency:** While the system functions well for smaller-scale experiments, larger and more complex tasks require further optimization. Addressing these computational efficiency challenges will be crucial for broader application.
- **Potential for Cross-Domain Application:** The framework has broad applicability across various scientific domains. Future work will be needed to customize the system for specific fields, ensuring that it meets the unique requirements of different areas of research.
- **Importance of Ethical Considerations:** Automated systems must be designed with strong ethical safeguards to prevent misuse or the generation of harmful research. This highlights the necessity of maintaining human oversight in automated processes.

5.5.2 Final Remarks

The success of this project underscores the significant impact that automation can have on accelerating scientific discovery, particularly for smaller research teams or institutions with limited resources. By providing a foundation for future advancements in automated research, this project aims to make scientific inquiry more efficient, accessible, and innovative.

References

- [1] Gemini Team et al. “Gemini: A Family of Highly Capable Multimodal Models”. In: *arXiv e-prints*, arXiv:2312.11805 (Dec. 2023), arXiv:2312.11805. DOI: 10.48550/arXiv.2312.11805. arXiv: 2312.11805 [cs.CL].
- [2] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. URL: <https://www.deeplearningbook.org/>.
- [3] Aaron Grattafiori et al. “The Llama 3 Herd of Models”. In: *arXiv e-prints*, arXiv:2407.21783 (July 2024), arXiv:2407.21783. DOI: 10.48550/arXiv.2407.21783. arXiv: 2407.21783 [cs.AI].
- [4] Chris Lu et al. “The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery”. In: *arXiv e-prints*, arXiv:2408.06292 (Aug. 2024), arXiv:2408.06292. DOI: 10.48550/arXiv.2408.06292. arXiv: 2408.06292 [cs.AI].
- [5] OpenAI. “GPT-4 Technical Report”. In: *arXiv preprint arXiv:2303.08774* (2023). Available online at <https://arxiv.org/abs/2303.08774>.
- [6] Adam Paszke, Sam Gross, Francisco Massa, et al. “PyTorch: An Imperative Style, High-Performance Deep Learning Library”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. Vol. 32. 2019. URL: https://papers.nips.cc/paper_files/paper/2019/hash/9015a7baeddc67fa85f0f9ab443eb861-Abstract.html.
- [7] Betty Shea and Mark Schmidt. “Why Line Search When You Can Plane Search? So-friendly Neural Networks Allow Per-iteration Optimization of Learning and Momentum Rates for Every Layer”. In: *arXiv abs/2406.17954* (2024). URL: <https://arxiv.org/abs/2406.17954>.
- [8] Qwen Team. *Qwen2.5: A Party of Foundation Models*. Sept. 2024. URL: <https://qwenlm.github.io/blog/qwen2.5/>.
- [9] Hugo Touvron et al. “LLaMA: Open and Efficient Foundation Models”. In: *Proceedings of the 2023 Conference on Natural Language Processing*. Available online at <https://arxiv.org/abs/2302.13971>. 2023.